

Cloud Computing Project

Lo-okup

Hemanth Aditya
Omid Askarisichani

Motivation

- An open source web application to add any picture to any place on Google Maps
- A social networks of users with the ability of like, rate, endorse and comment for each picture
- Be able to run on any device, and for the most part, compatible with all phone Operating Systems
- Highly available, consistent and fast

Application

Lookup allows users to pin places around them on a map for the rest of the community to explore.

Users can browse through places around them within a 50 miles radius. They can rate, review and share places.

Lookup

Lookup

Lookup

Login

Lookup

hemanth@g.com


Logout


Isla Vista, CA, United States

SEARCH

CARD VIEW

MAP VIEW


 **IV Deli Market** 14mi




0 reviews 3 ratings

Rate this place

★★★★☆


 **Santa Barbara Beautiful Museum** 14mi




0 reviews 3 ratings

Rate this place

★★★★☆

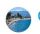
 **Labyrinth** 14mi




0 reviews 3 ratings

Rate this place


★★★★☆

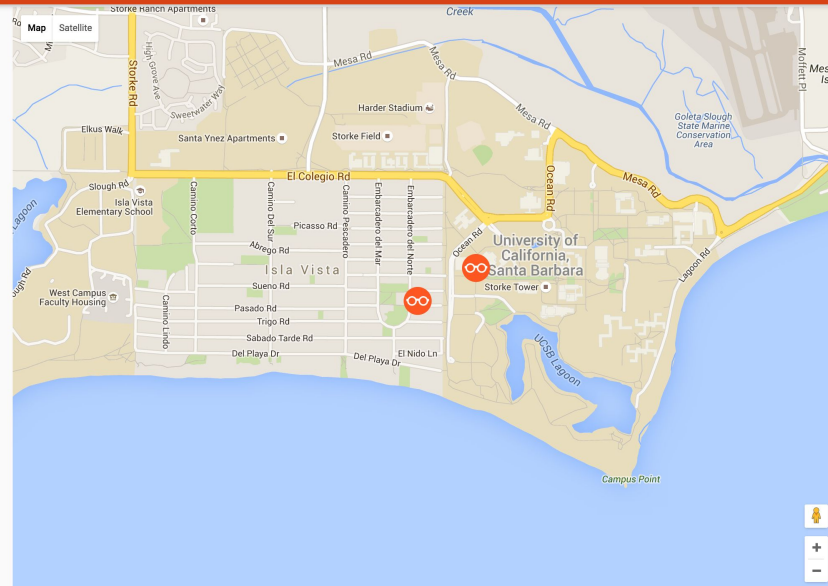

 **Goleta Beach** 14mi



 **Woodstock Pizza** 14mi



 **The Luxor** 14mi



Demo

Application: Framework

Ruby on Rails serves as the MVC framework of the application.

Notable gems:

- Geokit / Geokit-rails
- Devise
- Ratyrate
- Paperclip
- Unicorn
- nginx

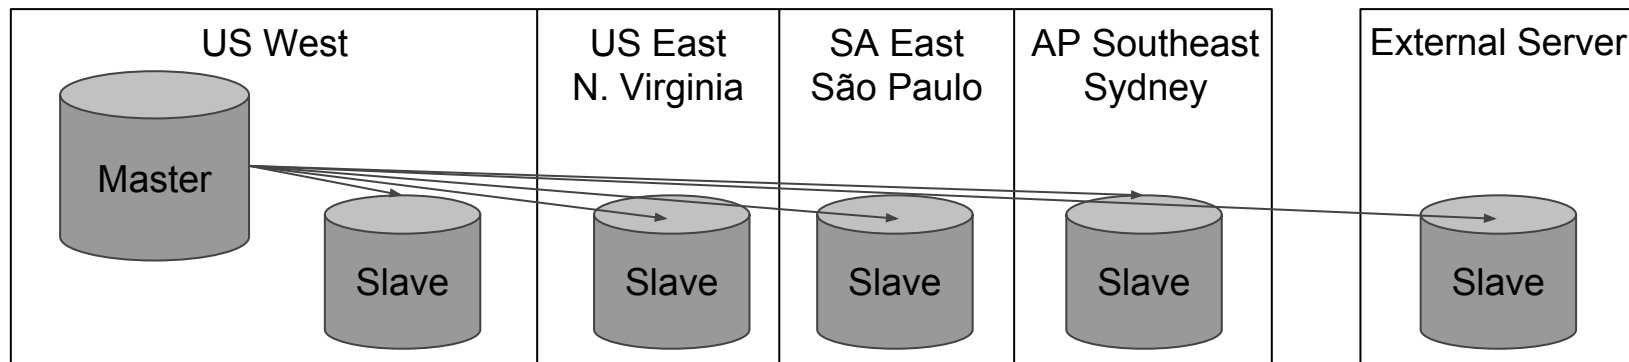
Cloud Infrastructure: Storage

- Data (coordinates, names, ratings and other information) are stored in the database backend
- Pictures (large entities) are stored in AWS S3 (and Google bucket storage)
- We use MySQL database because our application requires relational structure that SQL-based databases make more sense

Cloud Infrastructure: Database

In order to make the data available we use replication.

Master-Slave replication architecture on multiple regions:



Writes go to the master and reads go to the slaves.

Cloud Infrastructure: Deploy the Application

Application is written in Ruby on Rails framework and is accessible from Github.

<https://github.com/adityahemanth/lookup>

We deploy the app with Unicorn and Nginx on Ubuntu 14.04 instances:


- Unicorn is an application server that enables Rails application to process requests concurrently
- Since Unicorn is not designed to be accessed by users directly, we are using Nginx as a reverse proxy that will buffer requests and responses between users and application

Cloud Infrastructure: Deployment Details

- We created an instance, installed all prerequisites and built lookup app on it
- Then we created a bootable unicorn service as `/etc/init.d/unicorn_lookup`
- This service starts every time an instance is booted, serves lookup app in production environment and listens on `shared/sockets/unicorn.sock`
- Therefore, we run a Nginx service to redirect all requests on port 80 to the aforementioned unicorn socket
- Last but not least, we have to open firewall for any IP on port 80

Cloud Infrastructure: Scale in/out

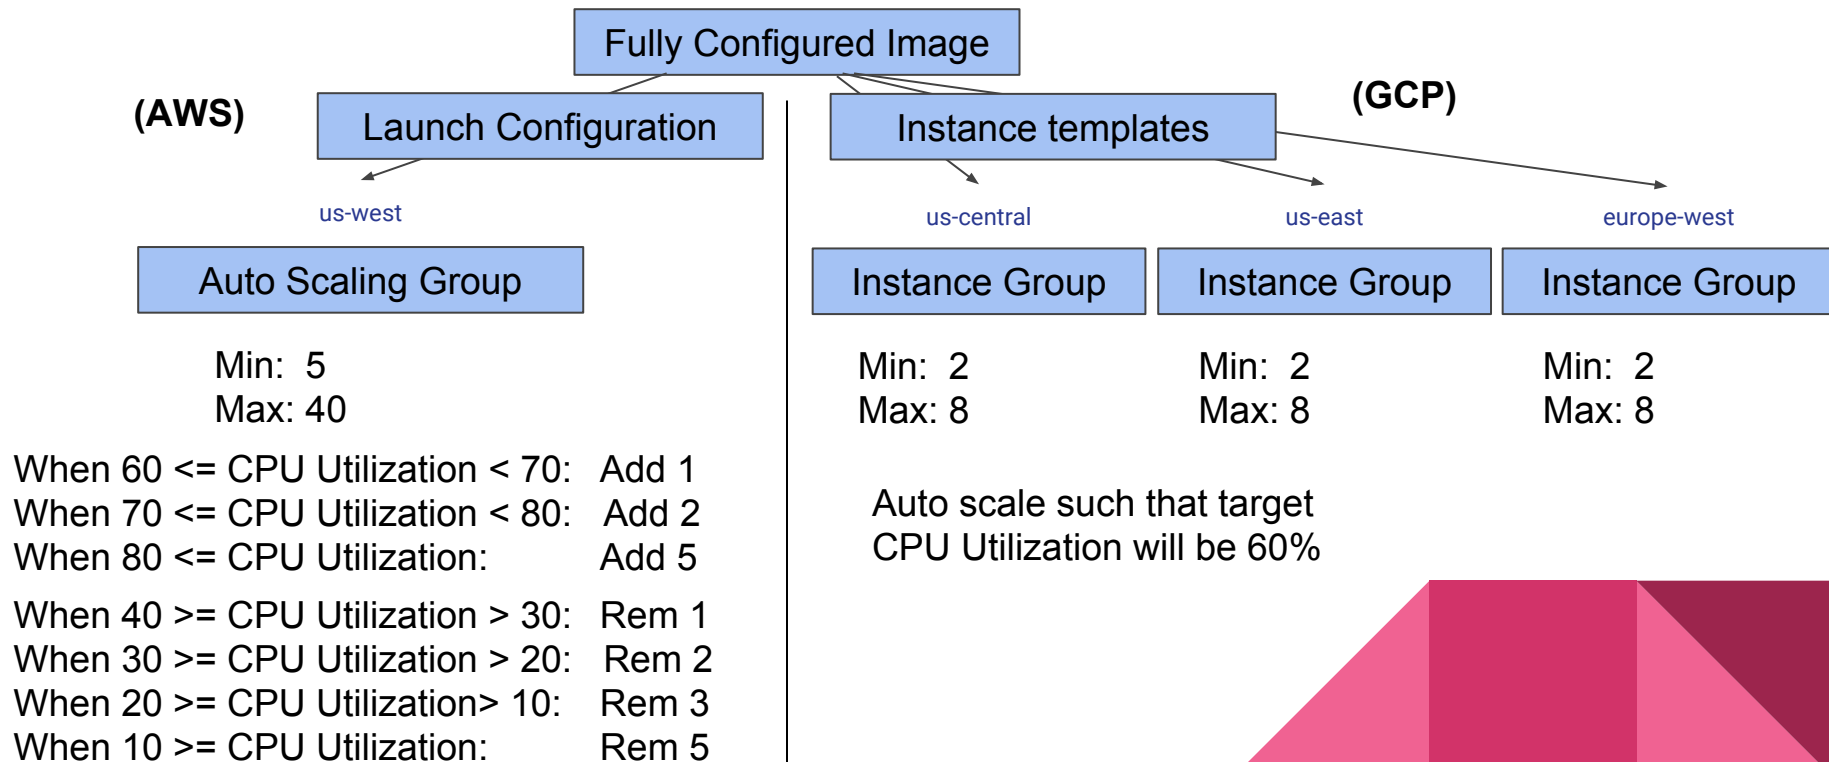
- Application is stateless and each instance can handle multiple users



- Thus, we need load balancers and auto scalars to do the rest for us


We are using Google Cloud Platform (GCP) and Amazon Web Services (AWS)

Cloud Infrastructure: Scale in/out




Cloud Infrastructure: Load Balancer

(AWS)



lookuplb-1945113182.us-west-1.elb.amazonaws.com


(GCP)



130.211.22.162

Cloud Infrastructure: DNS

- Route 53 for Round Robin DNS
- A weighted policy for two backend servers
- We use load balancers' addresses for DNS



Cloud Infrastructure: DNS

Traffic flow

Traffic policies

Policy records

Domains

Registered domains

Pending requests

Start point

DNS type



Weighted rule

Weight

Health checks

☒ Evaluate target health

lookuphc (320eeb17-9cf3-4bb1-9525-380c3a983121)

Weight

Health checks

☒ Evaluate target health

lookuphc (320eeb17-9cf3-4bb1-9525-380c3a983121)



Endpoint

Type

Value



Endpoint

Type

Value

```
C:\Windows\system32\nslookup.exe
*** Request to ResNet-0-3.resnet.ucsb.edu timed-out
> lo-okup.com
Server: ResNet-0-3.resnet.ucsb.edu
Address: 169.231.0.3

Non-authoritative answer:
Name: lo-okup.com
Address: 130.211.22.162

> lo-okup.com
Server: ResNet-0-3.resnet.ucsb.edu
Address: 169.231.0.3

DNS request timed out.
timeout was 2 seconds.
*** Request to ResNet-0-3.resnet.ucsb.edu timed-out
> lo-okup.com
Server: ResNet-0-3.resnet.ucsb.edu
Address: 169.231.0.3

Non-authoritative answer:
Name: lo-okup.com
Address: 130.211.22.162

> lo-okup.com
Server: ResNet-0-3.resnet.ucsb.edu
Address: 169.231.0.3

Non-authoritative answer:
Name: lo-okup.com
Addresses: 52.9.210.221
          54.241.161.140

> lo-okup.com
Server: ResNet-0-3.resnet.ucsb.edu
Address: 169.231.0.3

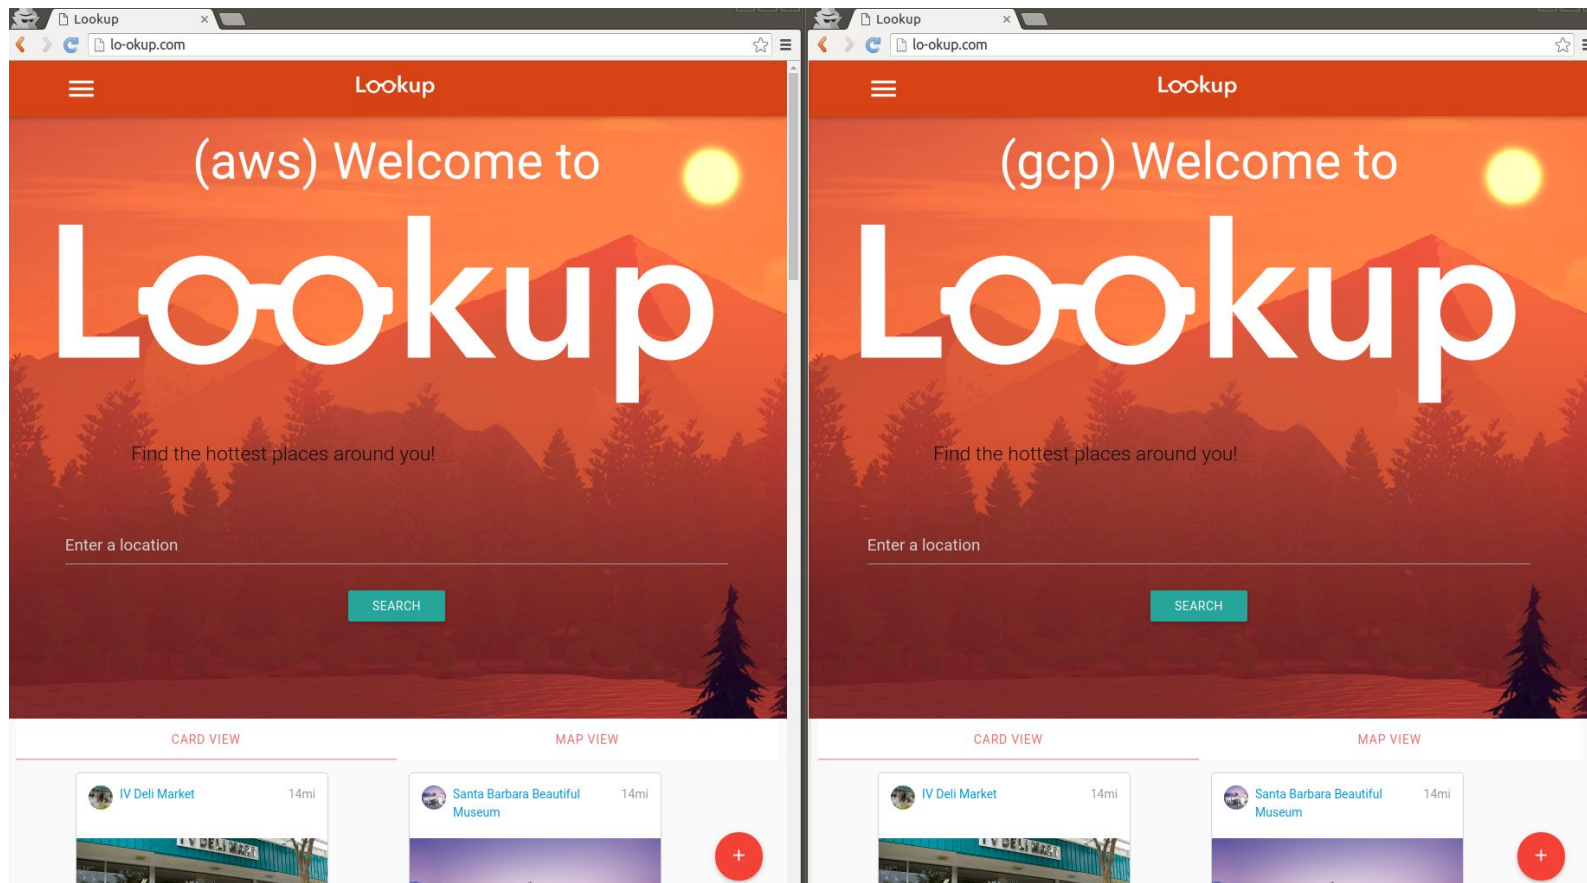
Non-authoritative answer:
Name: lo-okup.com
Addresses: 54.241.161.140
          52.9.210.221
```

All name systems

In order to have test the system properly, we created separate name systems (sub domains) as follows:


- s1.lo-okup.com
- aws.lo-okup.com
- gcp.lo-okup.com
- lo-okup.com

lo-okup.com



Test

- Using simple python script:



<input type="checkbox"/> Name ^	Zone	Network	In use by	Internal IP	External IP	Connect
<input type="checkbox"/> lookup-instance-group-2fun	europe-west1-b	default		10.132.0.2	146.148.28.67	SSH
<input type="checkbox"/> lookup-instance-group-3xky	europe-west1-b	default		10.132.0.4	130.211.64.106	SSH
<input checked="" type="checkbox"/> lookup-instance-group-bmxj	europe-west1-b	default	lookup-instance-group	10.132.0.3	104.155.23.12	SSH
<input checked="" type="checkbox"/> lookup-instance-group-e6qy	europe-west1-b	default	lookup-instance-group	10.132.0.5	104.155.21.85	SSH
<input checked="" type="checkbox"/> lookup-instance-group-s1fn	europe-west1-b	default	lookup-instance-group	10.132.0.7	146.148.9.31	SSH

```
In [ ]: # Omid55
        # Scalability check for servers
        import urllib.request as req
        import threading
```

```
In [ ]: def worker():
        lookupurl = 'http://gcp.lo-okup.com/'
        opens = 10000

        for i in range ( 1 , opens + 1 ):
            print ( i , end = " " )
            req.urlopen ( lookupurl )

        print( "DONE" )
```


```
In [ ]: threadsCnt = 5
        threads = []
        for i in range( threadsCnt + 1 ):
            t = threading.Thread( target = worker )
            threads.append( t )
            t.start()
```

```
In [ ]:
```

Test using Tsung


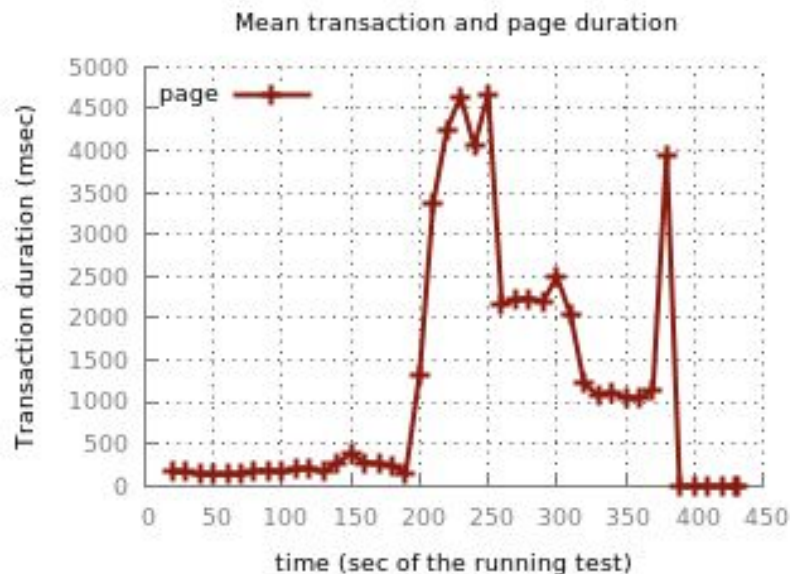
- We did not stop there for checking the availability and scaling properties!
- We used Tsung for HTTP load stress test

We start with 4 and ends
with 64 users in
incrementing order of time:



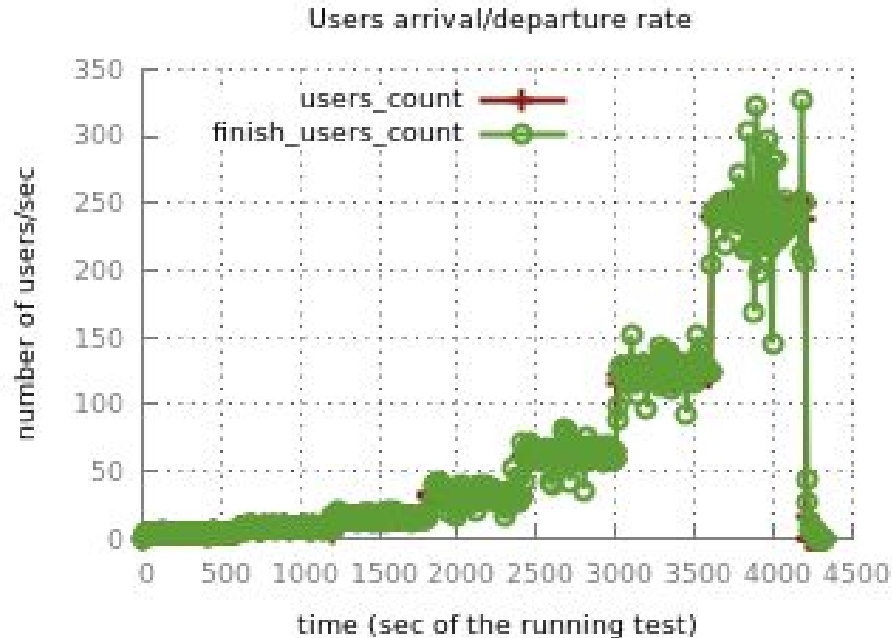
Test

- Test on s1.lo-okup.com shows one micro-cpu can handle 16 users/second:





Test

- We double number of users per second every 10 minutes
- It starts with 4 users/second and ends with 256 users/second



Test

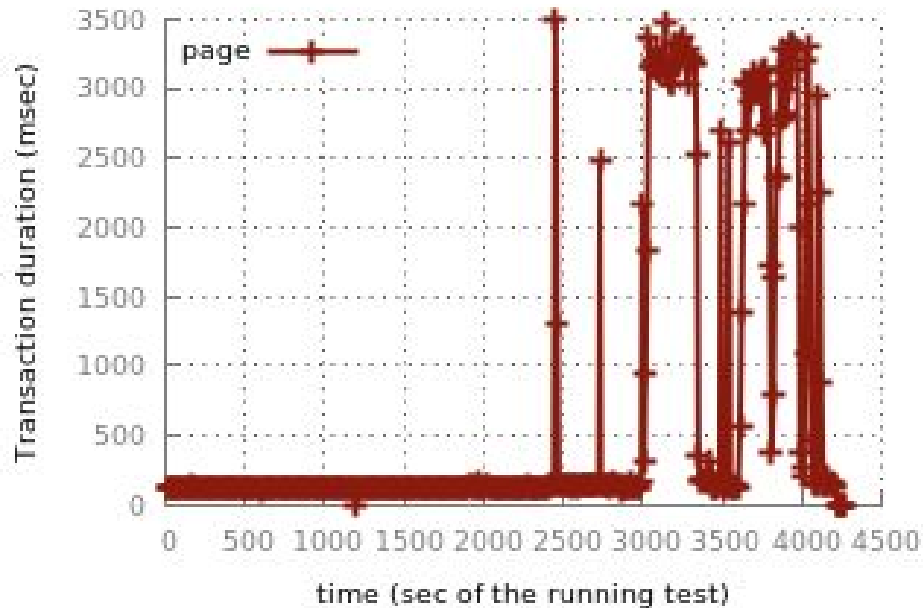
- Test on one of the load balancers (gcp.io-okup.com):




Test

- Test on aws system (aws.lo-okup.com):

Mean transaction and page duration




HTTP Code Response rate



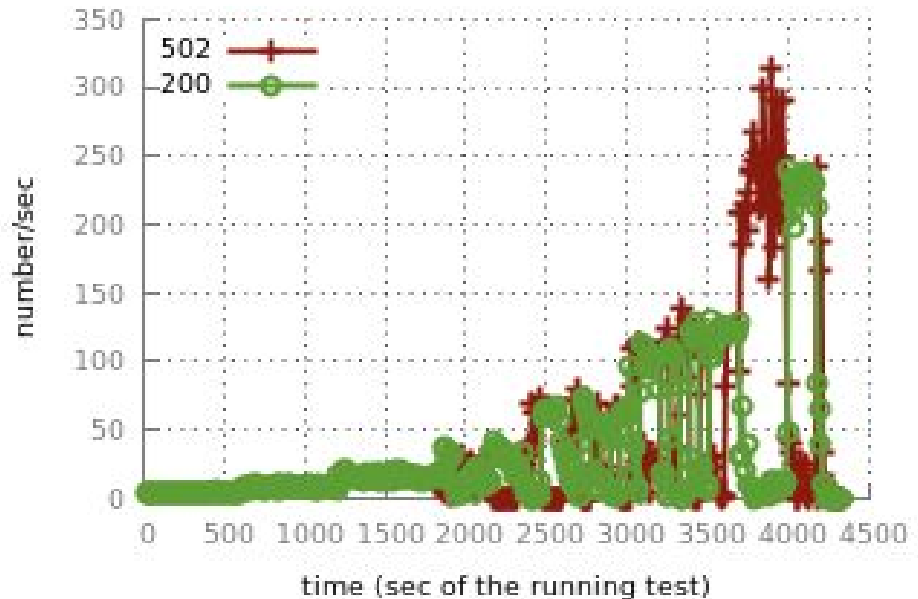
Test

- Test on whole system (lo-okup.com):

Mean transaction and page duration



HTTP Code Response rate



Future works

Route 53 dynamic weight changes.

Cost optimizations / load time optimizations

Thank you